

STAT 663

Statistical Graphics and Data Exploration I

Week 13

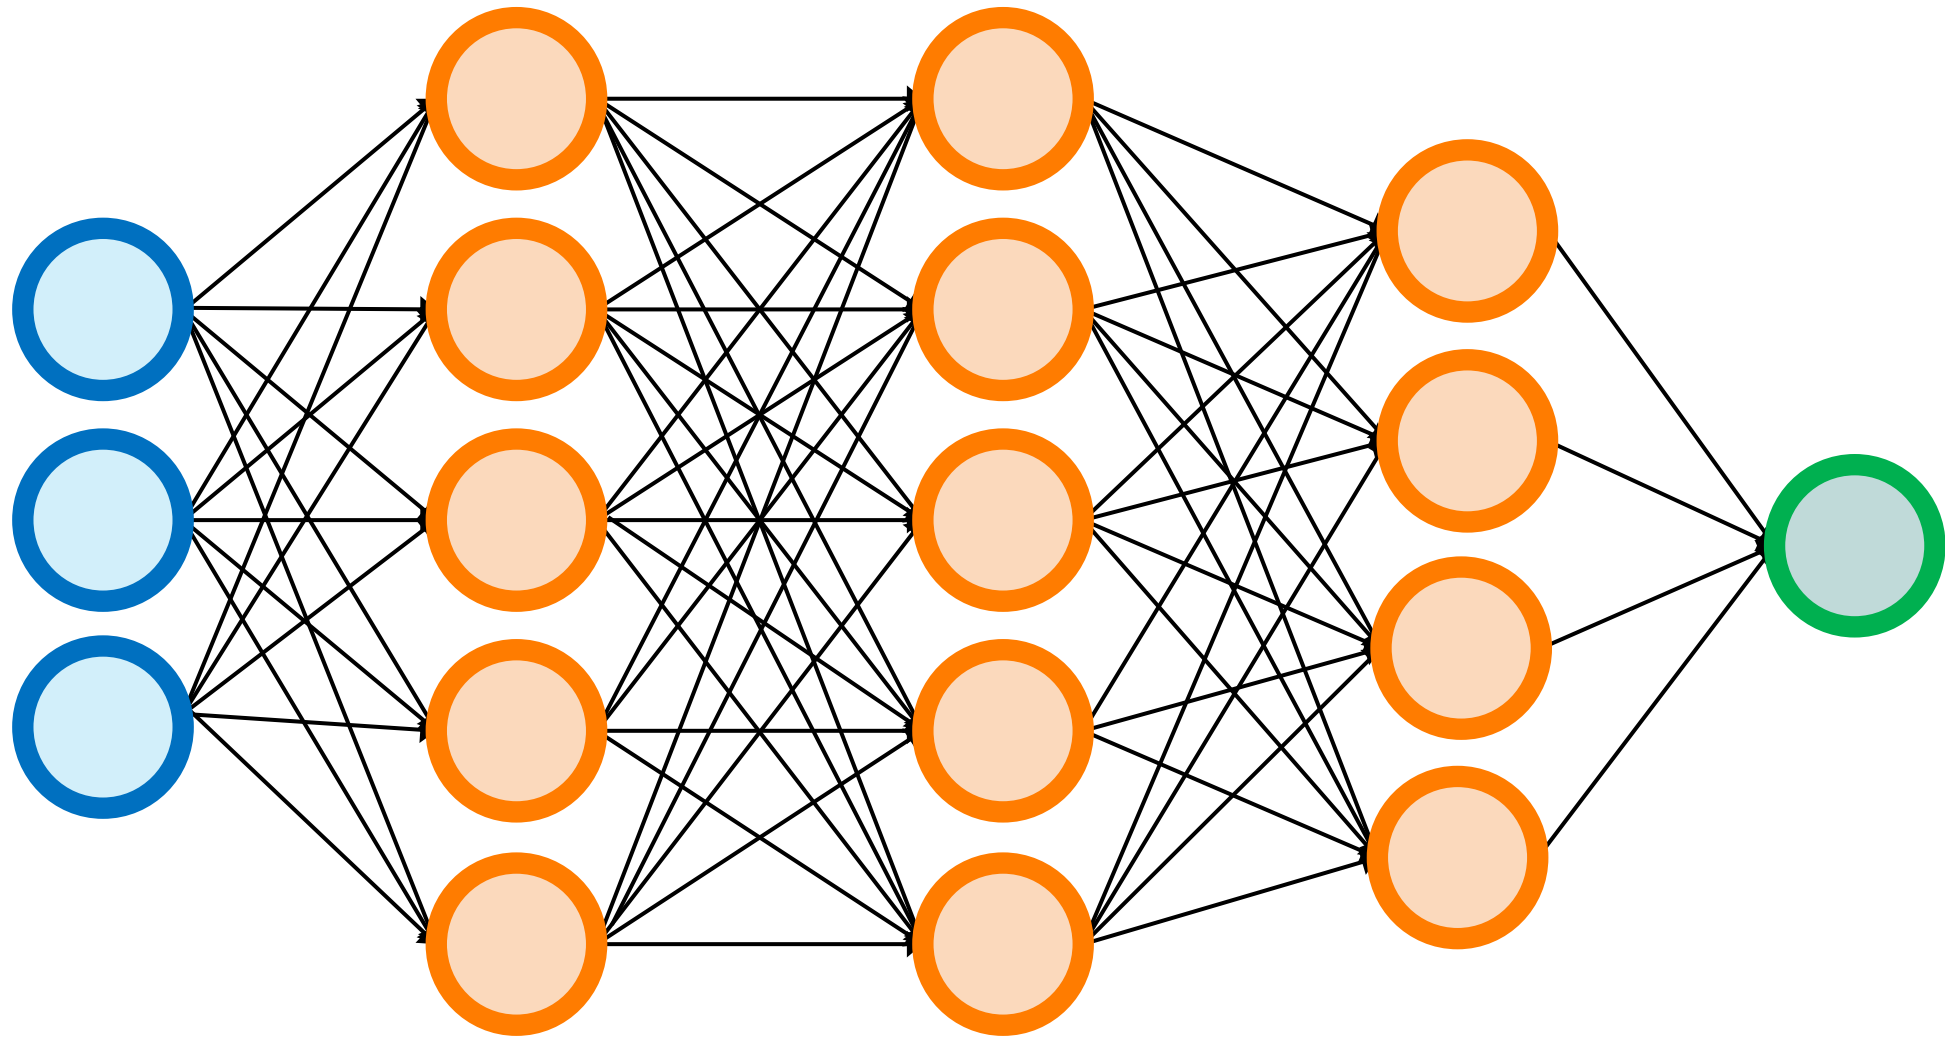
Neural Network

Lecture 2

Lesson Objectives

- Neural Network Auto-Regressive (NNAR) models
- The `nnetar()` function in R
- Application of the NNAR model to COVID-19 study

Neural Networks



Input Layer

Hidden Layers

Output Layer

Neural Network Autoregression (NNAR) Models

- Lagged values of the time series can be used as inputs to a neural network.
- $\text{NNAR}(p, k)$: p lagged inputs and k nodes in the single hidden layer.
- $\text{NNAR}(p, 0)$ model is equivalent to an $\text{ARIMA}(p, 0, 0)$ model but without stationarity restrictions.
- Seasonal $\text{NNAR}(p, P, k)_m$: inputs $(Y_{t-1}, Y_{t-2}, \dots, Y_{t-p}, Y_{t-m}, Y_{t-2m}, Y_{t-Pm})$ and k neurons in the hidden layer.
- $\text{NNAR}(p, P, 0)_m$ model is equivalent to an $\text{ARIMA}(p, 0, 0)(P, 0, 0)_m$ model but without stationarity restrictions.


NNAR models in R

- The `nnetar()` function in R can be used to fit an $\text{NNAR}(p, P, k)_m$ model.
- If p and P are not specified, they are automatically selected.
- For non-seasonal time series, default p = optimal number of lags (according to the AIC) for a linear $\text{AR}(p)$ model.
- For seasonal time series, defaults are $P = 1$ and p is chosen from the optimal linear model fitted to the seasonally adjusted data.
- Default $k = (p + P + 1)/2$ (rounded to the nearest integer).

NNAR Models: Rolling Forecast

- The network is applied iteratively.
- For forecasting one step ahead, we simply use the available historical inputs.
- For forecasting two steps ahead, we use the one-step forecast as an input, along with the historical data.
- This process proceeds until we have computed all the required forecasts.

Time Series Plot for Florida



The nnetar function

We can also specify the options in the `nnetar` function.

- The **repeats** option of `nnetar` shows **the number of networks** (default = 20) to fit with different random starting weights. These networks are then averaged when producing forecasts.
- The **size** option provides number of nodes in the hidden layer. Default is half of the number of input nodes (including external regressors, if given) plus 1.
- We can consider a Box-Cox transformation for the data.
 - If **lambda = auto**, then a transformation is automatically selected using `BoxCox.alpha`.
 - The transformation is ignored if **lambda = NULL**.
 - Otherwise, data transformed before model is estimated.

A Time Series Plot for Florida

```
library(slid)
library(dplyr)
data(state.ts)
Florida.ts <- state.ts %>%
dplyr::filter(State == "Florida")
y <- ts(Florida.ts$Y.Death, start = c(2020,1),
frequency = 7)
```

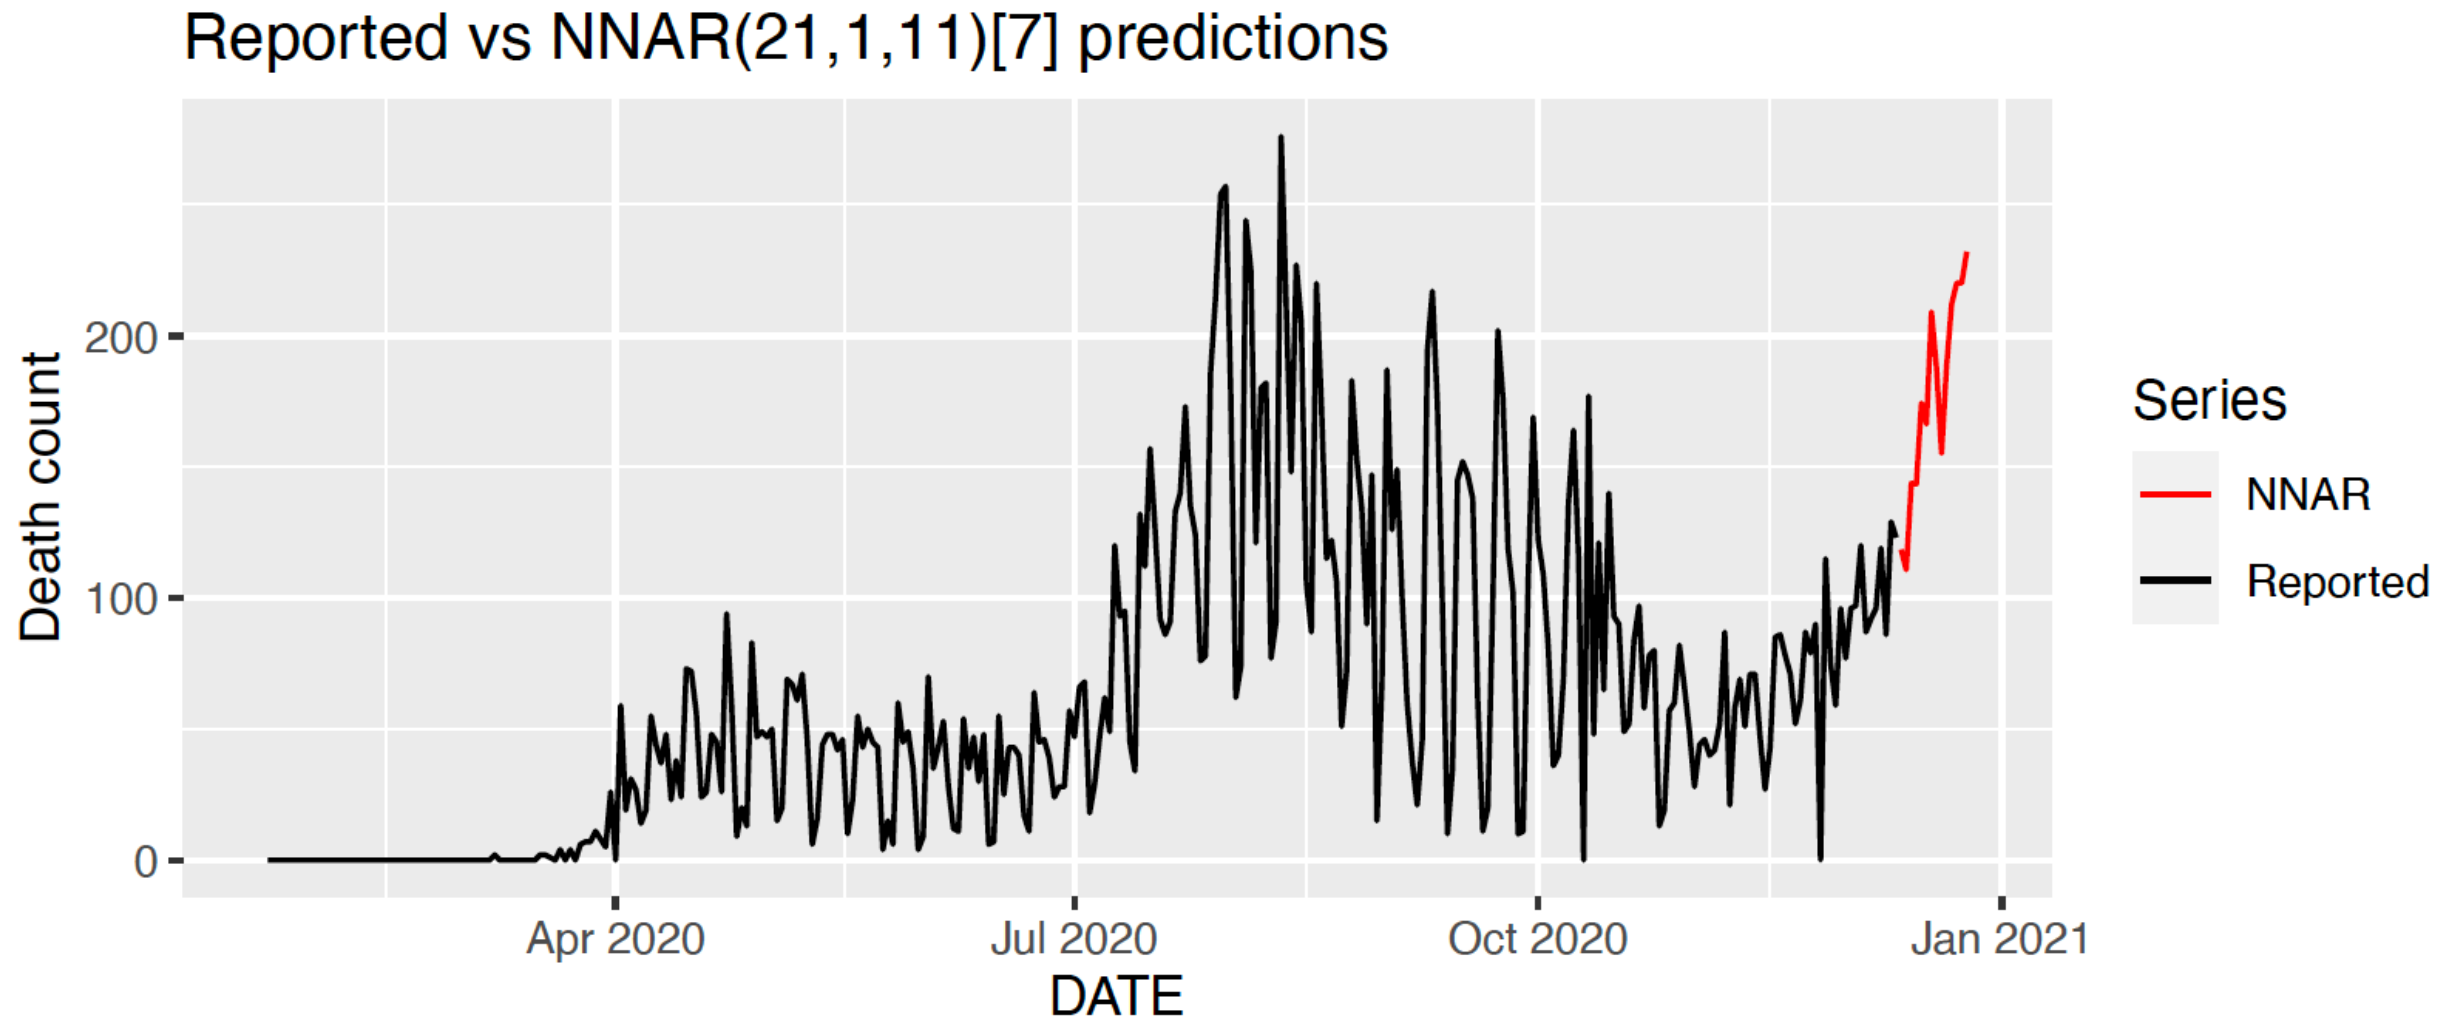
```
library(forecast)
set.seed(2020)
fit <- nnetar(y)

fit
## Series: y
## Model: NNAR(21,1,11)[7]
## Call: nnetar(y = y)
##
## Average of 20 networks, each of which is
## a 21-11-1 network with 254 weights
## options were - linear output units
##
## sigma^2 estimated as 6.598
```

```
nnar.fc.report <- as.data.frame(forecast(fit, h = 14)) %>%
mutate(DATE = as.Date("2020-12-11") + 1:14)
names(nnar.fc.report) <- c("Y.Death", "DATE")

# plot of reported vs NNAR predictions
ggplot() +
geom_line(aes(x = DATE, y = Y.Death, color = "Reported"),
data = Florida.ts) +
geom_line(aes(x = DATE, y = Y.Death, color = "NNAR"),
data = nnar.fc.report) +
scale_color_manual(
values = c(Reported = "black", NNAR = "red")) +
labs(y = "Death count",
title = "Reported vs NNAR(21,1,11)[7] predictions") +
guides(color = guide_legend(title = "Series"))
```

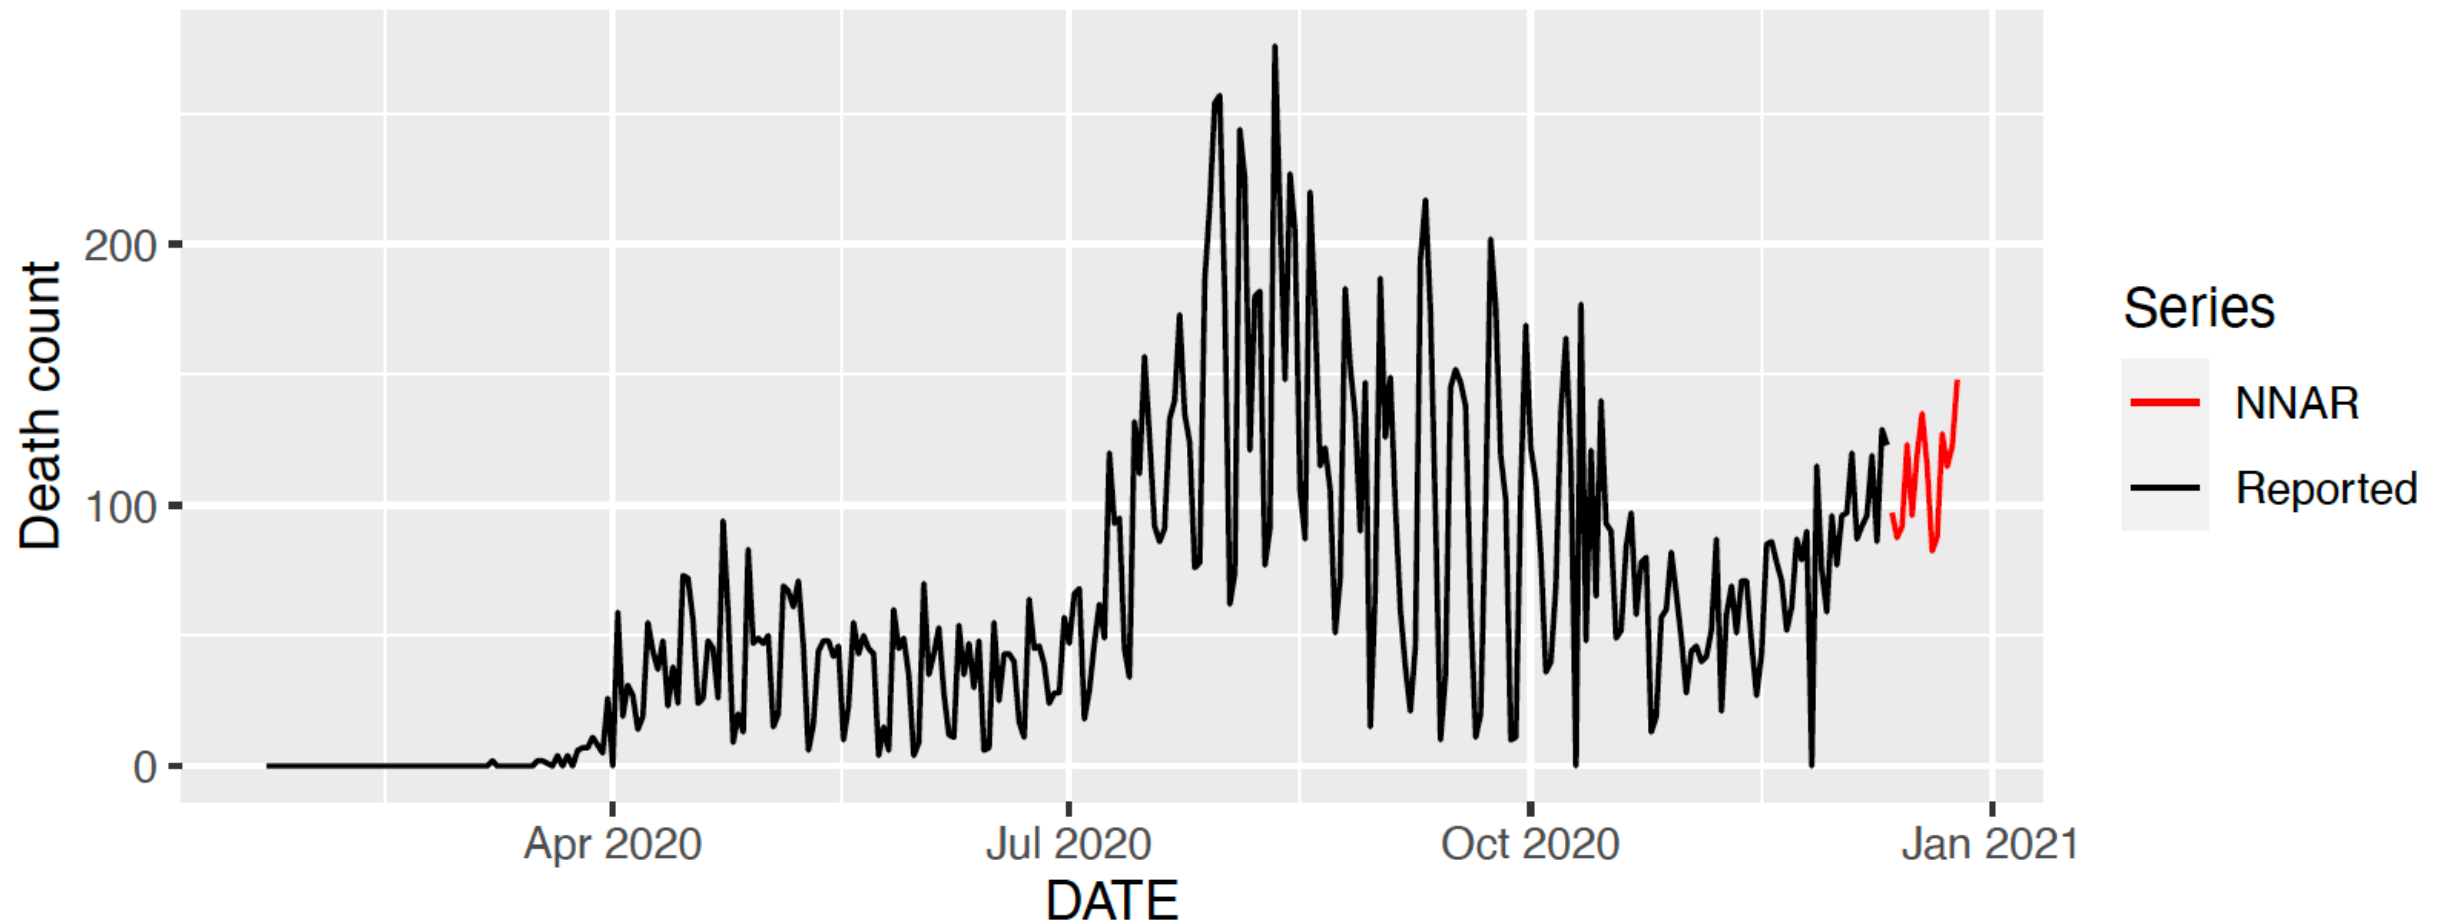
NNAR Prediction



Log Transformation

```
fit <- nnetar(y + 1, lambda = 0, size = 5)
```

Reported vs NNAR(14,1,5)[7] predictions



Prediction Intervals

- Define the NNAR(p, k) model as

$$Y_t = f(Y_{t-1}, \dots, Y_{t-p}) + \varepsilon_t,$$

where f is a neural network with k hidden nodes in a single layer. The error series $\{\varepsilon_t\}$ is assumed to be homogeneous.

- We can simulate future sample paths of this model iteratively, by randomly generating a value for ε_t , either from a normal distribution, or by resampling from the historical values.
 - So if ε_{T+1}^* is a random draw from the distribution of errors at time $T + 1$, then

$$Y_{T+1}^* = f(Y_T, \dots, Y_{T-p+1}) + \varepsilon_{T+1}^*$$

is one possible draw from the forecast distribution for Y_{T+1} .

- We can then repeat the process to get

$$Y_{T+2}^* = f(Y_{T+1}^*, Y_T, \dots, Y_{T-p+1}) + \varepsilon_{T+2}^*$$

and obtain a sample path $\{Y_{T+1}^*, Y_{T+2}^*, \dots, Y_{T+h}^*\}$.

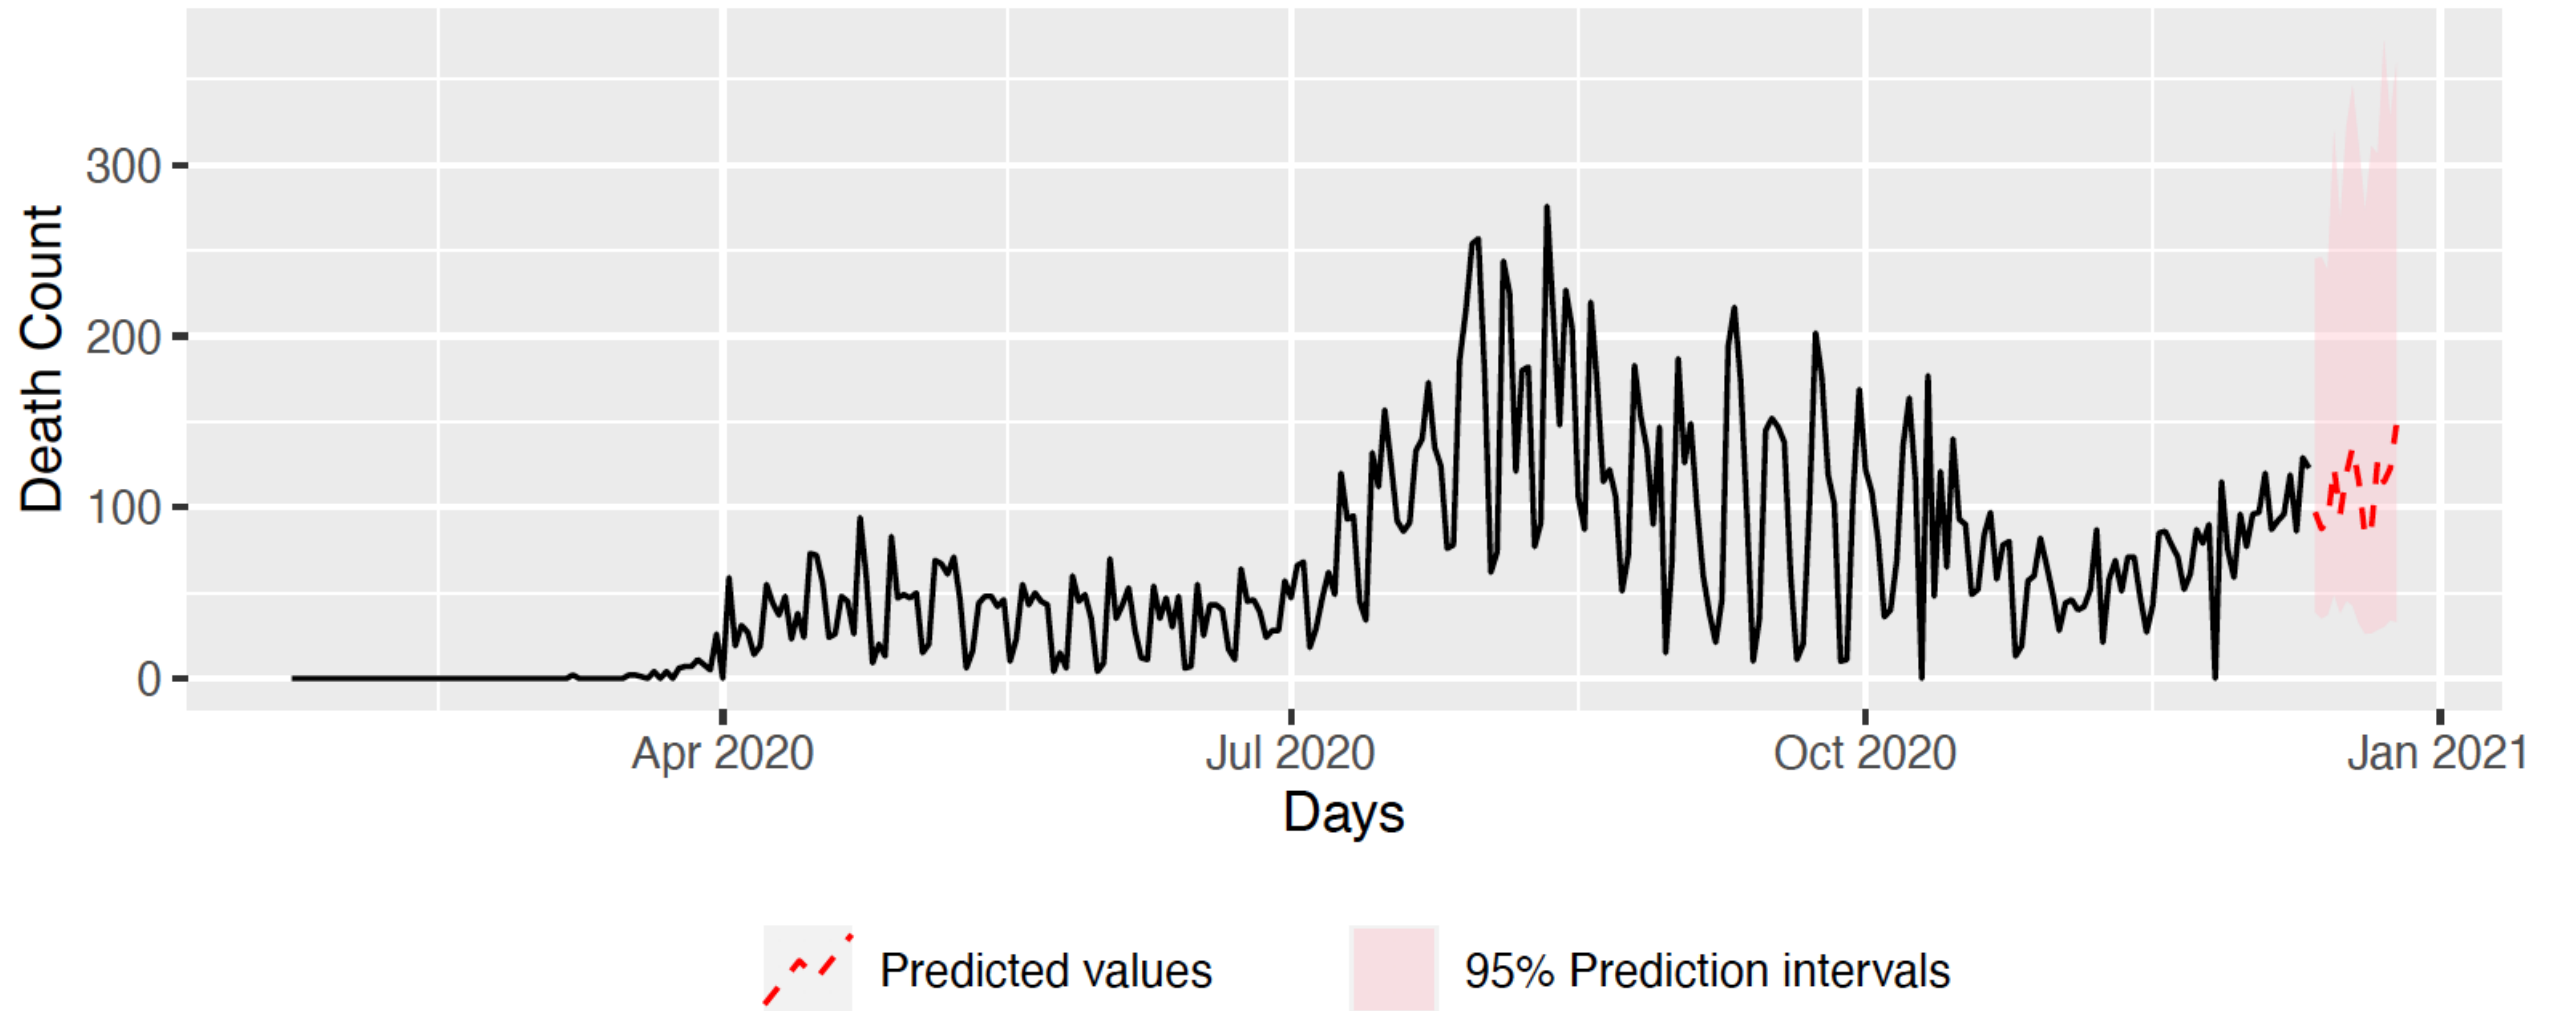
- If we generate many sample paths, we can get a good picture of the forecast distributions.

Prediction Intervals

```
fcast <- forecast(fit, PI = TRUE, h = 14, npaths = 500)
```

```
#autoplot(fcast, xlab = "Days", ylab = "Daily new deaths")
fcast.report <- as.data.frame(fcast) %>%
mutate(DATE = as.Date("2020-12-11") + 1:14)
names(fcast.report) <- c("Y.Death", "Lo.80", "Hi.80", "Lo.95", "Hi.95", "DATE")
ggplot(Florida.ts, aes(DATE, Y.Death)) +
geom_line() +
labs(x = "Days", y = "Death Count") +
# Add prediction intervals
geom_ribbon(mapping = aes(x = DATE, y = Y.Death, ymin = Lo.95, ymax = Hi.95,
fill = '95% Prediction intervals'), data = fcast.report, alpha = 0.4) +
# Add line for predicted values
geom_line(mapping = aes(x = DATE, y = Y.Death, colour = 'Predicted values'),
linetype = "dashed", data = fcast.report, key_glyph = "timeseries") +
scale_colour_manual("", values = "red") +
scale_fill_manual("", values = "pink") +
theme(legend.position = "bottom")
```

Prediction Intervals



Two Statistical Cultures

